

Practical Machine Learning Assignment

Chun-Fu Wang

Loading the Data

```
library(dplyr)
```

```
##  
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':  
##  
##   filter, lag
```

```
## The following objects are masked from 'package:base':  
##  
##   intersect, setdiff, setequal, union
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
trainData <- read.csv("pml-training.csv")  
testData <- read.csv("pml-testing.csv")  
  
trainData <- trainData[2:length(trainData)]  
testData <- testData[2:length(testData)]  
  
dim(trainData)
```

```
## [1] 19622 159
```

```
dim(testData)
```

```
## [1] 20 159
```

There are 19622 observations in the training data set and 20 observations in the test data set that we are going to predict.

Exploratory Data Analysis

Removing the near zero variance features as well as the statistically insignificant features.

```
# cleaning up data
dim(trainData)
```

```
## [1] 19622  159
```

```
nzv <- nearZeroVar(trainData)
filteredTrainData <- trainData[, -nzv]
filteredTestData <- testData[, -nzv]

# removed statistically insignificant variables
filteredTrainData <-
  filteredTrainData %>%
  select(-c(user_name,
             raw_timestamp_part_1,
             raw_timestamp_part_2,
             cvtd_timestamp,
             max_roll_belt:var_yaw_belt,
             var_accel_arm,
             max_pitch_arm:amplitude_yaw_arm,
             max_roll_dumbbell:amplitude_pitch_dumbbell,
             var_accel_dumbbell:var_yaw_dumbbell,
             max_pitch_forearm:amplitude_pitch_forearm,
             var_accel_forearm))

# remove from test set as well
filteredTestData <-
  filteredTestData %>%
  select(-c(user_name,
             raw_timestamp_part_1,
             raw_timestamp_part_2,
             cvtd_timestamp,
             max_roll_belt:var_yaw_belt,
             var_accel_arm,
             max_pitch_arm:amplitude_yaw_arm,
             max_roll_dumbbell:amplitude_pitch_dumbbell,
             var_accel_dumbbell:var_yaw_dumbbell,
             max_pitch_forearm:amplitude_pitch_forearm,
             var_accel_forearm))

dim(filteredTrainData)
```

```
## [1] 19622  54
```

```
dim(filteredTestData)
```

```
## [1] 20 54
```

Preprocess

Split the training data into two set of 80% and 20%.

```
set.seed(142678)

dataIndex <- createDataPartition(filteredTrainData$classe, p = 0.8, list = FALSE)
trainSet <- filteredTrainData[dataIndex, ]
testSet <- filteredTrainData[-dataIndex, ]
```

Machine Learning

Using Random Forest and Rpart to train the model.

```
library(doMC)
```

```
## Loading required package: foreach
```

```
## Loading required package: iterators
```

```
## Loading required package: parallel
```

```
registerDoMC(cores = 4)
```

```
modelRf <- train(classe ~ ., data = trainSet, model = "rf")
```

```
## Loading required package: randomForest
```

```
## randomForest 4.6-12
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':
##
##     margin
```

```
## The following object is masked from 'package:dplyr':
##
##     combine
```

```
modelRpart <- train(classe ~ ., data = trainSet, model = "rpart")

predRf <- predict(modelRf, newdata = testSet)
predRpart <- predict(modelRpart, newdata = testSet)

C1 <- confusionMatrix(predRf, testSet$classe)
print(C1)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1115    1    0    0    0
##           B    0  758    2    0    0
##           C    0    0  682    2    0
##           D    0    0    0  640    0
##           E    1    0    0    1  721
##
## Overall Statistics
##
##           Accuracy : 0.9982
##           95% CI : (0.9963, 0.9993)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9977
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9991  0.9987  0.9971  0.9953  1.0000
## Specificity      0.9996  0.9994  0.9994  1.0000  0.9994
## Pos Pred Value   0.9991  0.9974  0.9971  1.0000  0.9972
## Neg Pred Value   0.9996  0.9997  0.9994  0.9991  1.0000
## Prevalence       0.2845  0.1935  0.1744  0.1639  0.1838
## Detection Rate   0.2842  0.1932  0.1738  0.1631  0.1838
## Detection Prevalence 0.2845  0.1937  0.1744  0.1631  0.1843
## Balanced Accuracy 0.9994  0.9990  0.9982  0.9977  0.9997
```

```
C2 <- confusionMatrix(predRpart, testSet$classe)
print(C2)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1115    1    0    0    0
##           B    0  758    2    0    0
##           C    0    0  682    2    0
##           D    0    0    0  640    1
##           E    1    0    0    1  720
##
## Overall Statistics
##
##           Accuracy : 0.998
##           95% CI : (0.996, 0.9991)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9974
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity           0.9991  0.9987  0.9971  0.9953  0.9986
## Specificity           0.9996  0.9994  0.9994  0.9997  0.9994
## Pos Pred Value         0.9991  0.9974  0.9971  0.9984  0.9972
## Neg Pred Value         0.9996  0.9997  0.9994  0.9991  0.9997
## Prevalence             0.2845  0.1935  0.1744  0.1639  0.1838
## Detection Rate         0.2842  0.1932  0.1738  0.1631  0.1835
## Detection Prevalence   0.2845  0.1937  0.1744  0.1634  0.1840
## Balanced Accuracy       0.9994  0.9990  0.9982  0.9975  0.9990
```

Out of Sample Error

Out of sample error for Random Forest is:

```
print((1 - C1$overall[1]) * 100)
```

```
## Accuracy
## 0.1784349
```

Out of sample error for Decision Tree is:

```
print((1 - C2$overall[1]) * 100)
```

```
## Accuracy
## 0.2039256
```

Predictin Result

Predicting the result using Random Forest models because of the lower out of sample error.

```
predResult <- predict(modelRf, newdata = filteredTestData)
print(predResult)
```

```
## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```